

Introduction to Computer Networking

Transport

TCP Service Model

- Transmission Control Protocol
- used by 95% of Internet applications
- reliable, end-to-end, bidirectional byte-stream service
- Transport layer
- two way communication channel between TCP peers = connection
- At both ends TCP keeps state machine
- 3 way handshake
 - Establish connection
 - SYN = synchronize
 - A wants to establish connection with TCP layer at B
 - A sends base number it will use to identify bytes in the byte stream
 - “0” = numbers will start at zero
 - “1000” = numbers will start at thousand
 - SYN + ACK = synchronize, acknowledgment
 - B sends ACK because it is acknowledging A’s request, agreeing to establish connection
 - ACK
 - A wants to accept the connection
- Stream of bytes
- data in order
- Stream of bytes is delivered by TCP segments
 - A puts bytes from the stream into a TCP segment

- A hands it to the IP layer
- IP layer delivers it to B
- in practice TCP layer may need to be transmitted multiple times
 - Segment is dropped, A doesn't receive acknowledgement
- it isn't very efficient to send lots of data at once
- A and B finish sending data
 - close connection = "teardown"
 - tell each other they are closing connection, clean up the state associated with the state machine
- Teardown
 - FIN = finish
 - ACK = acknowledgement
 - This closes down the data stream from A to B
 - if B has new data to send to A, B can keep sending data
 - The ACK can also carry new data
 - FIN
 - ACK
- TCP service model
- Sequence numbers
 - every segments header carries the sequence number - in the stream of bytes - of the first byte in segment
 - Both sides agree that the sequence numbers start at 1000
 - first segment will have sequence number of 1000
 - segment carries 500 bytes then next segment will carry the sequence number 1500
 - segment gets lost -> sequence number will be incorrect

- if data arrives out of order -> resequenced them to the correct order using sequence number
- Flow control
 - if A is much faster than B
 - Receiver keeps telling the sender if it can keep sending
 - How much room has it for new data
- Congestion
 - TCP tries to divide up the network capacity equally among all the TCP connections using the network
- TCP segment header is much more complicated than IP and Ethernet headers
 - Because TCP connection is reliable
- Acknowledgment sequence number
 - Which byte we are expecting next
 - 501 - we received byte 500
- sequence numbers for every direction
- TCP options fields
 - Extra fields that were added after the TCP standard was created
 - Header Length field tells us how many option fields are present
- Flags
 - ACK
 - SYN
 - FIN
 - PSH
 - tells TCP layer at other end to deliver data immediately upon arrival
 - don't wait for more data
 - short segments carrying time critical data (key stroke)

- URG
- TCP connection is uniquely identified by 5 pieces of information in TCP and IP headers
 - IP source and destination address
 - IP Protocol ID for TCP - tells us the connection is TCP
 - TCP source and destination ports
- A must have unique source port ID
 - it can't use the same port it is already using with another connection to the same service on B
 - A minimize chances
 - Increments the source port number for every new connection
- ISN = initial sequence number
 - Random

UDP Service Model

- User Datagram Protocol
- no guarantee delivery service
 - Application doesn't need it
 - Application handles retransmission in its own way
- UDP takes application data -> create UDP datagram -> hands it to the network layer
- Four header fields
- Length
 - Data + header
- UDP checksum
 - Optional when using IPv4
 - sender doesn't include checksum - field is filled with 0s
 - UDP header + data

- Includes part of the IP header
 - allows the UDP layer to detect datagrams that were delivered to the wrong destination
- port numbers
- UDP demultiplexing Protocol
 - Demultiplexing mechanism to divide to the stream of UDP datagrams and send them to the correct process
- connectionless datagram service
 - no connection is establish because none is needed (no handshake)
 - all of the information is self contained in the datagram
 - any order of packet - application needs to resequenced the data itself
- unreliable delivery
 - no acknowledgement that data reached the other end
 - cant detect missing datagrams
 - application can ask for the data again by itself
 - build retransmission mechanism on top of UDP
 - early version of NFS (network file system) did this
 - they didn't want to use the sliding window used by TCP
 - they created their own inside application
- application that doesn't need reliable delivery
 - simple request-response application
 - DNS
 - domain name system
 - turns a hostname into an IP address
 - uses UDP
 - request is fully contained in one UDP datagram

- DHCP
 - uses UDP for same reasons
- real-time streaming audio
 - less common to use UDP instead of TCP today

ICMP Service Model

- Internet Control Message Protocol
- used to report errors and diagnose problems with the network layer
- 3 mechanisms that we use to make the network layer work in the internet
 - IP
 - creation of IP datagrams
 - Hop by hop delivery from end to end
 - Routing tables
 - ICMP
 - Communicates network layer information between end hosts and routers
 - reports error conditions
 - Help diagnose problems
- runs above network layer, transport layer protocol
- end host or router want to report an error using ICMP
 - it puts information into an ICMP payload
 - hands it to IP -> send
- Example
 - Web client running as application
 - Client wants to access an HTTP server
 - packet comes to router
 - Address that is put in here is to a network that this router has no information about in its forwarding table

- router will send back a message “destination network unreachable”
- ICMP doesn't attempt to resend its message,
- Doesn't maintain any state
- router wants to send back ICMP message
 - it takes the header
 - Takes source and destination address
 - Place it into an ICMP message
 - Type of error and its code
- ICMP message types
 - Destination Network Unreachable
 - Router doesn't know network where it should send packet
 - Destination Host Unreachable
 - IP datagram gets to the last router but then the last router doesn't know where the host is
 - Destination Port Unreachable
 - it doesn't recognize port or protocol ID at the other end
- Ping command
 - Uses ICMP
 - Test the liveness and connectivity of another host
 - ping 156.1.1.1 or ping www.seznam.cz
 - Calls ICMP directly, sends ICMP echo request
 - B echo reply
- Traceroute
 - application that tells us the path that packets take through the network
 - Traceroute seznam.cz
 - path and round trip delay to each of the routers

- ICMP
- A sends UDP message to B
 - Time to live in IP header is set to one
 - router must discard the message when TTL is 0 and send back ICMP message
- Router sends ICMP message to A “TTL expired”
- and so on until datagram gets to B
- B sends back “port unreachable”
 - A chose weird port that B will not know

End-to-End Principle

- Refers to 2 principles
 - Correctness
 - Strong end-to-end principle
- why doesn't network layer do more
 - Compress data, add security, reformat
- end to end principle
 - Network can do many things to help
 - end points must be responsible
 - Nobody else has the information necessary to do this correctly
 - you want your application secure = end-to-end security implemented in the application
- End points also have to detect errors, not just network
 - link error detection was designed for errors in transmission, not errors in storage
- Wire link layers are very reliable but wireless are not
- Stronger than the first one
 - First one said that you have to implement something end-to-end but you can also implement it in the middle for performance improvements

- strong principle says not to implement it in the middle, only in the end
 - Flexibility and simplicity

Error Detection

- Network error detection algorithms
 - Checksum
 - cyclic redundancy code (CRC)
 - Message authentication code (MAC)
- We have a payload of data
 - we calculate some error detection bits over that data
 - either append it or prepend it to the payload
 - Ethernet - appends cyclic redundancy code
 - Transport Layer Security - appends a message authentication code
 - IP - prepends a checksum placed in IP header
 - TLS and Ethernet - footers, protocol information which follow the payload
- 3 commonly used error detection algorithms
 - Checksum
 - IP, TCP
 - easy to compute
 - weak error detection guarantees
 - Cyclic Redundancy code
 - Ethernet
 - Harder to compute
 - computes the remain of a polynomial
 - Message Authentication Code
 - TLS
 - Cryptographic transformation of data

- Combines the packet with some secret information to generate a value
- Someone can only generate or check the MAC if they have the secret
- so if you receive a packet and its MAC is correct then you are sure that computed MAC has the secret
- https
- not great for catching errors, mainly for security

Finite State Machines

- finite state machine is composed of a finite number of states
- State
 - Particular configuration of the system
- Edges between states define how we transition between them
 - What events cause transition to occur
 - what actions the system will take when that transition occurs
 - not all transition have actions associated with them
 - system is in state and an event arrives for which there is no transition -> FSM is undefined
 - Multiple transition from single state
- Example - HTTP request
 - Page = idle
 - we want to load a new page requesting state
 - event is load a new page
 - Action is open a connection to the web server
 - Request pending state
 - Having more resources to request, requesting recourse with HTTP GET
- if you want to be explicit, what happens on each state for every event
 - What happens if the connection closes when we are in request pending state

- TCP connection
 - 12 state
 - 4 parts
 - Top 4 states
 - Describes how you open a TCP connection
 - Center state “establish”
 - TCP is sending and receiving data
 - Bottom 6 states
 - How connection close
 - Bottom state “close”
 - Connection is closed and node can forget about it

Reliable Communications - Stop and Wait

- Flow control
- Basic problem
 - sender can send data faster than receiver can process it
 - Slower processor, networking card
- Approach
 - Don't send more packets than B can process
 - Receiver gives feedback
 - stop and wait
 - sliding window
- Principles
 - At most one packet in flight at any time
 - Waits for acknowledgement
 - If it doesn't receive packet -> assume that packet is lost -> resend
 - sender sends one packet

- 4 cases
 - nothing is lost, data is lost, acknowledgement is lost, acknowledgement is delayed
 - acknowledgement is delayed
 - the second ACK is sent by receiver -> sender doesn't know whether this ACK is for retransmission or new data packet -> we can have error
 - finite state machine has to keep track of it
- how do you detect duplicates?
 - Is ACK for retransmission or duplicated copies of packets vs new data
 - solution
 - one big counter on all data and ACK
 - Receiver can tell if new data or duplicate
 - simplifying assumptions
 - Network doesn't duplicate packets
 - Packets are not delayed multiple timeouts
 - these problems can be solved by increasing sequence number

Reliable Communication - Sliding Window

- Used in most high performance protocols
- Stop and wait
 - 10 Mbps link -> Boston node can receive 10 Mbps
 - Round trip time (RTT) is 50 milliseconds
 - We send Ethernet frames (Ethernet frame = max 12 kb)
 - $1000/50 \text{ milliseconds} = \text{we can send 20 packet a second}$
 - $20 \text{ packets} * 12 \text{ kb} = 240 \text{ kbps}$
 - Stop and wait protocol is using 2% of what could be used
 - Inefficient

- Sliding window
 - Generalization of stop and wait
 - Allow multiple un-ACKed segments
 - $n=5$
 - San Francisco can have 5 packets in flight
 - Boston can have 5 ACK on flight
 - Bound on number of un-ACKed segments called window
 - can keep pipe full
 - $10 \text{ Mbps} / 20 \text{ trip times} = 500 \text{ kb per trip time}$
 - $500 \text{ kb} / 12 \text{ kb} = 41 \text{ packets}$
 - it allows 10 Mbps connection
- sliding window sender
 - Every segment has sequence number
 - Sender maintains 3 variables
 - send window size (SWS)
 - Last acknowledgement received (LAR)
 - Last segment sent (LSS)
 - Maintain invariant ($LSS - LAR \leq SWS$)
 - Receive ACK \rightarrow advance LAR on new ACK \rightarrow buffer up to SWS segment
 - Stalling window
 - $n=3$, send packet 1, 2, 3
 - packet 1 is lost, 2 and 3 are ACK
 - Sender can't send packet 4,... until it receives ACK for packet 1
 - Although most of the data in the window has been delivered it can't move past the first unACKed piece of data
- sliding window receiver

- Receiver maintains 3 variables
 - Receive window size (RWS)
 - Last acceptable segment (LAS)
 - Last segment received (LSR)
- Maintains invariant ($LAS - LSR \leq RWS$)
- if received packet is $< LAS$, send ACK
 - Cumulative ACK
 - what is the end of contiguous data you received
 - if received 1,2,3,5 -> ACK 3
 - note TCP ACKs are next expected data (ACK 4)
 - $n+1$
 - first missing byte
- RWS, SWS and sequence space
 - $RWS \geq 1$, $SWS \geq 1$, $RWS \leq SWS$
 - if $RWS = 1$ "go back N" protocol, need $SWS+1$ sequence numbers
 - Example
 - $RWS = 1$, $SWS = 3$
 - A sends 0,1,2 -> B ACKs 0,1,2
 - when B ACK 0 -> A sends 3
 - When B ACK 1 -> A sends 4
 - A sends 3,4,5
 - 3 is dropped
 - B receives 4,5 -> sends ACK for them
 - A times out and resends 3
 - RWS was size one so receive can't buffer 4 and 5
 - sender has to retransmit 3,4,5

- if $RWS=SWS$, need $2SWS$ sequence numbers
- Generally need $RWS+SWS$ sequence numbers
 - RWS packets in unknown state (ACK doesn't have to arrive)
 - SWS packets in flight must not overflow sequence number space
- TCP flow control
 - TCP is sliding window protocol and uses flow control
 - Receiver advertises RWS using window field
 - sender can only send data up to $LAR + \text{window}$
 - Receiver specifies the flow control window using the window field
 - = buffer size, how much I can accept
 - TCP receiver will only handle data equal to the acknowledged sequence number plus the window
 - so sender isn't allowed to send data past $ACK + \text{window}$
- Example
 - $RWS=2, SWS=3$
 - A sends 0,1,2
 - 0 is ACK -> send 3, 1 is ACK -> send 4, 2 is ACK send 5
 - 3 is ACK -> send 6
 - 4 lost -> B waits for 4 -> send another ACK for 3 (cumulative ACK)
 - A resends 4 -> 4 arrives to B
 - $RWS=2$ so 5 is buffered, B can also ACK 5

Reliable Communications - Retransmission Strategies

- per-packet timer
- Strategies
 - Go back N
 - single is lost -> retransmit outstanding window of packets

- assumes all packets are lost
- selective repeat
 - Retransmit only lost packet
 - assumes only one packet is lost
- if there is bunch of losses, selective repeat protocol is slower
- Better to use go back n
 - Because receiver has RWS=1 sender will have to retransmit every single packet in the window
- you shouldn't retransmit earlier than you should
 - Inflating number of packets that are beyond your window size

Reliable Communications - TCP header

- 20 bytes (5x32 octets)
- Options after TCP header
- Acknowledgement number=5000 - received up to 5000 byte
 - Acknowledges last byte received + 1
- Sequence number=4000 and 500 bytes data - byte 4000 to 4499
 - Receiver sends ACK number=4500
- ACK packet
 - TCP segment that has no data
 - Traffic is unidirectional
 - ACK added to TCP segment with data
 - Traffic is bidirectional
- Checksum
 - Computed over TCP pseudo header
 - TCP header + some of the IP header + data
- Control bits - U, A, P, R, S, F

- U
 - Urgent bit
- P
 - Push bit
 - push data to the receiving application
- A = ACK
- R = reset
 - Reset connection
- S = SYN
 - “this is my starting sequence number, synchronize to this number”
- F = FIN
- it is better not to have starting sequence number 0
 - Security problems - people can guess it
- Offset
 - what offset within the segment does data begin

Reliable Communications Connection Setup and Teardown

- 3 way handshake
 - SYN, S_n
 - random sequence number
 - SYN, S_p, ACK, S(a+1)
 - S(a+1), ACK, S(p+1)
 - Sequence number - where bytes would start but there are no byte
- Simultaneous open
 - another way how to open connection
 - if both side know each other port numbers
 - Both sides send SYNs at the same time

- how to open connection
 - SYN, Sa
 - SYN, Sp
 - SYN, Sa, ACK, S(p+1)
 - SYN, Sp, ACK S(a+1)
- Wireshark
 - Sequence number = 0
 - It actually isn't 0
 - Wireshark display relative sequence number - relative to beginning of the stream
- connection teardown
 - FIN but says no more data to send
 - Caused by close() or shutdown() on other end
 - Both sides must send FIN to terminate a connection
 - Typical teardown exchange
 - A -> B: FIN, seq Sa, ack Sb
 - B -> A: ack S(a+1)
 - B -> A: FIN, seq Sb, ack S(a+1)
 - A -> B: ack S(b+1)
 - can also have simultaneous close
- when can we close connection?
 - Problems with closed socket
 - Final ack is lost
 - same port is immediately reuse for a new connection
 - Solution
 - Active closer goes into time wait before it can reuse its state
 - Keep socket around for 2 maximum segment lifetime

Monday, May 29, 2017

- this solution can cause problems to servers
 - OS has too many sockets in time wait
 - there are trick - send a reset, delete the socket, set socket option to time 0
 - OS won't let you restart server because port still in use
- for both - active and passive